

Advanced Machine Learning



Bogdan Alexe,

bogdan.alexe@fmi.unibuc.ro

University of Bucharest, 2nd semester, 2019-2020

Recap - Shattering

Definition (restriction of \mathcal{H} to C)

Let \mathcal{H} be a set hypothesis, i.e., set of functions from \mathcal{X} to $\{0, 1\}$, and let C be a (finite) subset of \mathcal{X} , $C = \{c_1, c_2, \dots, c_m\}$. The restriction of \mathcal{H} to C , denoted by \mathcal{H}_C , is the set of functions from C to $\{0, 1\}$ that can be derived from \mathcal{H} . That is:

$$\mathcal{H}_C = \{h: C \rightarrow \{0, 1\} \mid h \in \mathcal{H}\}$$

Definition (Shattering)

A hypothesis class \mathcal{H} *shatters* a finite set C of \mathcal{X} , if the restriction of \mathcal{H} to C is the set of all functions from C to $\{0, 1\}$. That is $|\mathcal{H}_C| = 2^{|C|}$.

Recap - The VC-dimension

Definition (VC-dimension)

The VC - dimension of a hypothesis class \mathcal{H} , denoted $\text{VCdim}(\mathcal{H})$, is the maximal size of a set $C \subseteq X$ that can be shattered by \mathcal{H} . If \mathcal{H} can shatter sets of arbitrarily large size we say that \mathcal{H} has infinite VC-dimension.

In order to show that the VC-dimension of a hypothesis class \mathcal{H} is d , we need to show that:


1. There exists a set C of size d that is shattered by \mathcal{H} . ($\text{VCdim}(\mathcal{H}) \geq d$)
2. Every set C of size $d + 1$ is not shattered by \mathcal{H} . ($\text{VCdim}(\mathcal{H}) < d+1$)

We will see in the next lecture that the converse is also true: *a finite VC-dimension guarantees learnability. Hence, the VC-dimension characterizes PAC learnability. VC-dimension is a combinatorial measure, does not imply computing probabilities.*


Recap - $\text{VCdim}(\mathcal{H}_{\text{thresholds}})$

Consider $\mathcal{H} = \mathcal{H}_{\text{thresholds}}$ be the set of threshold functions over the real line.


$$\mathcal{H}_{\text{thresholds}} = \{h_a: \mathbb{R} \rightarrow \{0, 1\}, h_a(x) = \mathbf{1}_{[x < a]}, a \in \mathbb{R}\}, |\mathcal{H}_{\text{thresholds}}| = \infty.$$



Consider $C = \{c_1\}$. Then $\mathcal{H}_C = \{h: C \rightarrow \{0, 1\} \mid h \in \mathcal{H}\}$ has two elements $\{h_a, h_b\}$ with $a \leq c_1$ and $b > c_1$ so \mathcal{H} shatters C . $\mathcal{H}_C = \{(0), (1)\}$, $|\mathcal{H}_C| = 2^{|C|} = 2^1$



Consider $C = \{c_1, c_2 \mid c_1 \leq c_2\}$. Then $\mathcal{H}_C = \{h: C \rightarrow \{0, 1\} \mid h \in \mathcal{H}\}$ has at most three elements, there is no function that realizes the labeling (0,1) and so \mathcal{H} does not shatter C .



So, **$\text{VCdim}(\mathcal{H}_{\text{thresholds}}) = 1$**


Recap - VCdim($\mathcal{H}_{\text{intervals}}$)

Consider $\mathcal{H} = \mathcal{H}_{\text{intervals}}$ be the set of intervals over the real line.


$$\mathcal{H}_{\text{intervals}} = \{[a,b] \mid a \leq b, a, b \in \mathbf{R}\}$$

Can also view $\mathcal{H}_{\text{intervals}}$ as:


$$\mathcal{H}_{\text{intervals}} = \{h_{a,b}: \mathbf{R} \rightarrow \{0, 1\}, h_{a,b} = \mathbf{1}_{[a,b]}, a \leq b, a, b \in \mathbf{R}\}$$



$\mathcal{H}_{\text{intervals}}$ shatters any set A of two different points in \mathbf{R} .



$\mathcal{H}_{\text{intervals}}$ cannot shatter any set A of three points in \mathbf{R} .



So, **VCdim($\mathcal{H}_{\text{intervals}}$) = 2**

$\text{VCdim}(\mathcal{H}_{\text{lines}}), \text{VCdim}(\mathcal{H}_{\text{rec}}^2)$

Consider $\mathcal{H} = \mathcal{H}_{\text{lines}}$ be the set of lines in \mathbf{R}^2 .

$$\mathcal{H}_{\text{lines}} = \{h_{a,b,c}: \mathbf{R}^2 \rightarrow \{0, 1\}, h_{a,b,c}((x,y)) = \mathbf{1}_{[ax+by+c>0]}((x,y)), a, b, c \in \mathbf{R}\}$$

$\mathcal{H}_{\text{lines}}$ shatters any set A of three non-collinear points in \mathbf{R}^2 .

$\mathcal{H}_{\text{lines}}$ doesn't shatter any set A of four points in \mathbf{R}^2 (geometric argument).

So, **$\text{VCdim}(\mathcal{H}_{\text{lines}}) = 3$**

Consider $\mathcal{H} = \mathcal{H}_{\text{rec}}^2$ be the set of axis aligned rectangles in the \mathbf{R}^2 .

$$\mathcal{H}_{\text{rec}}^2 = \{[a,b] \times [c,d] \mid a \leq b, c \leq d, a, b, c, d \in \mathbf{R}\}$$

$\mathcal{H}_{\text{rec}}^2$ shatters the set A of 4 points arranged as a diamond. So $\text{VCdim}(\mathcal{H}_{\text{rec}}^2) \geq 4$

$\mathcal{H}_{\text{rec}}^2$ doesn't shatter any set A of five points in \mathbf{R}^2 (geometric argument).

So, **$\text{VCdim}(\mathcal{H}_{\text{rec}}^2) = 4$**

Today's lecture: Overview

- Computing the VC-dimension for some particular \mathcal{H}
- Assignment 1

Some basic properties of the $\text{VCdim}(\mathcal{H})$

1. $\text{VCdim}(\mathcal{H}) \leq \log_2 |\mathcal{H}|$
2. If $\mathcal{H}_1 \subseteq \mathcal{H}_2$ then $\text{VCdim}(\mathcal{H}_1) \leq \text{VCdim}(\mathcal{H}_2)$
3. If $\text{VCdim}(\mathcal{H}) = \infty$ then \mathcal{H} is not PAC learnable

VCdim(\mathcal{H}_{\sin})


$$\text{VCdim}(\mathcal{H}_{\text{thresholds}}) = 1, \text{VCdim}(\mathcal{H}_{\text{intervals}}) = 2, \text{VCdim}(\mathcal{H}_{\text{lines}}) = 3$$
$$\text{VCdim}(\mathcal{H}_{\text{rec}}^2) = 4$$

Consider $\mathcal{H} = \mathcal{H}_{\sin}$ be the set of sin functions:

$$\mathcal{H}_{\sin} = \{h_{\theta}: \mathbf{R} \rightarrow \{0,1\} \mid h_{\theta}(x) = \lceil \sin(\theta x) \rceil, \theta \in \mathbf{R}\}, \lceil -1 \rceil = 0$$

$$\text{VCdim}(\mathcal{H}_{\sin}) = ?$$

We will show that $\text{VCdim}(\mathcal{H}_{\sin}) = \infty$



● Class 0

● Class 1

$\text{VCdim}(\mathcal{H}_{\sin})$

Lemma

Let $x \in (0, 1)$ and let $0.x_1x_2x_3\dots$ be the binary representation of x . Then, for any natural number m , provided that there exist $k \geq m$ such that $x_k = 1$, we have:

$$\left\lceil \sin(2^m \pi x) \right\rceil = 1 - x_m$$

Example of binary representation:

$$x = (0.x_1x_2x_3\dots)_2 = x_1 \times 2^{-1} + x_2 \times 2^{-2} + x_3 \times 2^{-3} + \dots$$

$$x = 0.75 = \frac{1}{2} + \frac{1}{4} = (0.110000\dots)_2$$

$$x = 0.3 = 0 \times 2^{-1} + 1 \times 2^{-2} + 0 \times 2^{-3} + 0 \times 2^{-4} + 1 \times 2^{-5} + \dots = (0.01001\dots)_2$$

$VCdim(\mathcal{H}_{\sin})$

Lemma

Let $x \in (0, 1)$ and let $0.x_1x_2x_3\dots$ be the binary representation of x . Then, for any natural number m , provided that there exist $k \geq m$ such that $x_k = 1$, we have:

$$\left\lceil \sin(2^m \pi x) \right\rceil = 1 - x_m$$

Proof

$$\begin{aligned} \sin(2^m \pi x) &= \sin(2^m \pi (0.x_1x_2x_3\dots)) = \sin(2\pi * 2^{m-1} (0.x_1x_2x_3\dots)) = \\ &= \sin(2\pi * (x_1x_2x_3\dots x_{m-1}.x_m x_{m+1}\dots)) \text{ (left shift with } m-1 \text{ position)} \\ &= \sin(2\pi * (x_1x_2x_3\dots x_{m-1}.x_m x_{m+1}\dots) - 2\pi * (x_1x_2x_3\dots x_{m-1}.0)) \text{ (sin has period } 2\pi) \\ &= \sin(2\pi * (0.x_mx_{m+1}\dots)) \end{aligned}$$

Note that $0.x_mx_{m+1}\dots > 0$ as there exist $k \geq m$ such that $x_k = 1$

Case 1: $x_m = 0$, then $0 < 2\pi * (0.x_mx_{m+1}\dots) < 2\pi * 1/2 = \pi$. So $0 < \sin(2^m \pi x) < 1$, and from here it results that: $\left\lceil \sin(2^m \pi x) \right\rceil = 1 = 1 - 0 = 1 - x_m$

Case 2: $x_m = 1$, then $2\pi > 2\pi * (0.x_mx_{m+1}\dots) \geq 2\pi * 1/2 = \pi$. So $-1 \leq \sin(2^m \pi x) \leq 0$, and from here it results that: $\left\lceil \sin(2^m \pi x) \right\rceil = 0 = 1 - 1 = 1 - x_m$ (we consider $\left\lceil -1 \right\rceil = 0$)

VCdim(\mathcal{H}_{sin})

To prove $\text{VCdim}(\mathcal{H}_{\text{sin}}) = \infty$, we need to pick n points which are shattered by \mathcal{H}_{sin} , for any n . To do so, we construct n points $x_1, x_2, \dots, x_n \in [0, 1]$, such that the set of the m -th bits in the binary expansion, as m ranges from 1 to 2^n , ranges over all possible labelings of x_1, x_2, \dots, x_n .

$$\begin{array}{rcl} x_1 & = & 0.0000 \dots 11 \\ x_2 & = & 0.0000 \dots 11 \\ & \dots & \\ x_{n-1} & = & 00011 \dots 11 \\ x_n & = & 0.0101 \dots 01 \end{array}$$

$m=1$

For example, to give the labeling 1 for all instances, we just pick $m=1$:

$$h(x) = \left\lceil \sin(2^1 \pi x) \right\rceil = 1 - \text{first_bit_of_binary_repres_of_}x$$

$$\mathcal{H}_{\text{sin}} = \{h_\theta: \mathbf{R} \rightarrow \{0,1\} \mid h_\theta(x) = \left\lceil \sin(\theta x) \right\rceil, \theta \in \mathbf{R}\},$$

which returns 1 - the first bit (column) in the binary expansion.

VCdim(\mathcal{H}_{\sin})

To prove $\text{VCdim}(\mathcal{H}_{\sin}) = \infty$, we need to pick n points which are shattered by \mathcal{H}_{\sin} , for any n . To do so, we construct n points $x_1, x_2, \dots, x_n \in [0, 1]$, such that the set of the m -th bits in the binary expansion, as m ranges from 1 to 2^n , ranges over all possible labelings of x_1, x_2, \dots, x_n .

$$\begin{array}{rcl} x_1 & = & 0.00000\dots 11 \\ x_2 & = & 0.00000\dots 11 \\ & \dots & \\ x_{n-1} & = & 0.0011\dots 11 \\ x_n & = & 0.0101\dots 01 \end{array}$$

$m=2$

If we wish to give the labeling 1 for x_1, x_2, \dots, x_{n-1} , and the labeling 0 for x_n , we pick $m=2$:

$$h(x) = \left\lceil \sin(2^2 \pi x) \right\rceil = 1 - \text{second_bit_of_binary_repres_of_}x$$

which returns 1 - the second bit (column) in the binary expansion.

VCdim(\mathcal{H}_{\sin})

To prove $\text{VCdim}(\mathcal{H}_{\sin}) = \infty$, we need to pick n points which are shattered by \mathcal{H}_{\sin} , for any n . To do so, we construct n points $x_1, x_2, \dots, x_n \in [0, 1]$, such that the set of the m -th bits in the binary expansion, as m ranges from 1 to 2^n , ranges over all possible labelings of x_1, x_2, \dots, x_n .

$$\begin{array}{rcl} x_1 & = & 0.0000\dots 11 \\ x_2 & = & 0.0000\dots 11 \\ & \dots & \\ x_{n-1} & = & 0.0011\dots 11 \\ x_n & = & 0.0101\dots 01 \end{array}$$

$m=2^n$

If we wish to give the labeling 0 for $x_1, x_2, \dots, x_{n-1}, x_n$ we pick $m=2^n$:

$$h(x) = \left\lceil \sin(2^{2^n} \pi x) \right\rceil$$

which returns 1 - the bit on position 2^n (column) in the binary expansion.

$\text{VCdim}(\mathcal{H}_{\text{sin}})$

To prove $\text{VCdim}(\mathcal{H}_{\text{sin}}) = \infty$, we need to pick n points which are shattered by \mathcal{H}_{sin} , for any n . To do so, we construct n points $x_1, x_2, \dots, x_n \in [0, 1]$, such that the set of the m -th bits in the binary expansion, as m ranges from 1 to 2^n , ranges over all possible labelings of x_1, x_2, \dots, x_n .

$$x_1 = 0.0000\dots 11$$

$$x_2 = 0.0000\dots 11$$

...


$$x_{n-1} = 0.0011\dots 11$$

$$x_n = 0.0101\dots 01$$

We conclude that $x_1, x_2, \dots, x_{n-1}, x_n$ can be given any labeling by some $h \in \mathcal{H}_{\text{sin}}$, so it is shattered. This can be done for any n , so $\text{VCdim}(\mathcal{H}_{\text{sin}}) = \infty$.

VCdim(\mathcal{H}_{\sin})

We conclude that $x_1, x_2, \dots, x_{n-1}, x_n$ can be given any labeling by some $h \in \mathcal{H}_{\sin}$, so it is shattered. This can be done for any n , so $\text{VCdim}(\mathcal{H}_{\sin}) = \infty$.



Training points


● Class 0

● Class 1

$$h(x) = \lceil \sin(20x) \rceil$$


$\text{VCdim}(\mathcal{H}_{\sin})$

We conclude that $x_1, x_2, \dots, x_{n-1}, x_n$ can be given any labeling by some $h \in \mathcal{H}_{\sin}$, so it is shattered. This can be done for any n , so $\text{VCdim}(\mathcal{H}_{\sin}) = \infty$.



- Class 0
- Class 1

$$h(x) = \lceil \sin(20x) \rceil$$




- Class 0
- Class 1

$$h(x) = \lceil \sin(77x) \rceil$$


$\text{VCdim}(\mathcal{H}_{\sin})$

We conclude that $x_1, x_2, \dots, x_{n-1}, x_n$ can be given any labeling by some $h \in \mathcal{H}_{\sin}$, so it is shattered. This can be done for any n , so $\text{VCdim}(\mathcal{H}_{\sin}) = \infty$.



- Class 0
- Class 1

$$h(x) = \lceil \sin(20x) \rceil$$




- Class 0
- Class 1

$$h(x) = \lceil \sin(77x) \rceil$$


$\text{VCdim}(\mathcal{H}_{\sin})$

We conclude that $x_1, x_2, \dots, x_{n-1}, x_n$ can be given any labeling by some $h \in \mathcal{H}_{\sin}$, so it is shattered. This can be done for any n , so $\text{VCdim}(\mathcal{H}_{\sin}) = \infty$.



- Class 0
- Class 1

$$h(x) = \lceil \sin(20x) \rceil$$



Label assigned

- Class 0
- Class 1

$$h(x) = \lceil \sin(77x) \rceil$$

VCdim($\mathcal{H}S_0^n$)

Consider $\mathcal{H} = \mathcal{H}S^n$ be the set of halfspaces (linear classifiers) in \mathbf{R}^n

$$\mathcal{H} = \mathcal{H}S^n = \{h_{w,b}: \mathbf{R}^n \rightarrow \{-1, 1\}, h_{w,b}(x) = \text{sign}\left(\sum_{i=1}^n w_i x_i + b\right) \mid w \in \mathbf{R}^n, b \in \mathbf{R}\}$$

Consider label -1 to correspond to label 0, they are basically the same.

For $n = 2$ we have:

$$\mathcal{H}S^2 = \mathcal{H}_{\text{lines}} = \{h_{a,b,c}: \mathbf{R}^2 \rightarrow \{0, 1\}, h_{a,b,c}((x,y)) = \mathbf{1}_{[ax+by+c>0]}((x,y)), a, b, c \in \mathbf{R}\}$$

Let us restrict our attention to “homogenous” linear classifiers, the ones that go through origin, $b = 0$.

$$\mathcal{H}S_0^n = \{h_{w,0}: \mathbf{R}^n \rightarrow \{-1, 1\}, h_{w,0}(x) = \text{sign}\left(\sum_{i=1}^n w_i x_i\right) \mid w \in \mathbf{R}^n\}$$

What is the VCdim($\mathcal{H}S_0^n$)?

VCdim(\mathcal{HS}_0^n)


$$\mathcal{HS}_0^n = \{h_{w,0}: \mathbf{R}^n \rightarrow \{-1, 1\}, h_{w,0}(x) = \text{sign}\left(\sum_{i=1}^n w_i x_i\right) \mid w \in \mathbf{R}^n\}$$

For $n = 2$ we have:

$$\mathcal{HS}_0^2 = \{h_{w_1, w_2}: \mathbf{R}^2 \rightarrow \{-1, 1\}, h_{w_1, w_2}(x) = \text{sign}(w_1 x_1 + w_2 x_2) \mid (w_1, w_2) \in \mathbf{R}^2\}$$

What is the VCdim(\mathcal{HS}_0^2) ?

\mathcal{HS}_0^2 shatters any set A of two different points.



Does \mathcal{HS}_0^2 shatter a set A of three points?

Difficult to reason geometrically... choose the algebraic proof.

VCdim(\mathcal{HS}_0^n)

We will show that $\text{VCdim}(\mathcal{HS}_0^n) = n$.

Proof: 1st part

We first show that $\text{VCdim}(\mathcal{HS}_0^n) \geq n$.

We find a set A consisting of n points in \mathbf{R}^n that is shattered by \mathcal{HS}_0^n .

Take $A = \{e_1, e_2, \dots, e_n\}$ to be the orthonormal basis of \mathbf{R}^n .

$e_1 = (1, 0, 0, \dots, 0)$; $e_2 = (0, 1, 0, \dots, 0)$; \dots ; $e_n = (0, 0, 0, \dots, 1)$

We want to proof that \mathcal{HS}_0^n shatters A , so that $\text{VCdim}(\mathcal{HS}_0^n) \geq n$. This is equivalent to proof that for every $B \subseteq A$, there is a function $h_B \in \mathcal{HS}_0^n$ such that h_B gives label +1 to all elements in B and label -1 to all elements of $A \setminus B$.

Pick B subset of A , $B \subseteq \{e_1, e_2, \dots, e_n\}$. Choose $w = (w_1, w_2, \dots, w_n)$ such that:

$$w_i = \begin{cases} 1, & \text{if } e_i \in B \\ -1, & \text{if } e_i \notin B \end{cases}$$

Then, $h_B(e_i) = \text{sign}(\langle w, e_i \rangle) = w_i$ will generate the labels +1 for elements in B , -1 for elements not in B

VCdim(\mathcal{HS}_0^n)

Proof: 2nd part

We now show that $\text{VCdim}(\mathcal{HS}_0^n) < n + 1$.

We will prove that given any set $A = \{x_1, x_2, \dots, x_{n+1}\}$ of $n + 1$ points in \mathbf{R}^n , A cannot be shattered by \mathcal{HS}_0^n .

The points $\{x_1, x_2, \dots, x_{n+1}\}$ “live” in \mathbf{R}^n , a vector space with dimension n . So, $\{x_1, x_2, \dots, x_{n+1}\}$ are linearly dependent and there exist coefficients a_1, a_2, \dots, a_{n+1} not all of them 0 such that:

$$\sum_{i=1}^{n+1} a_i x_i = 0$$

Take $P \subseteq \{1, 2, \dots, n+1\}$ the set of positive coefficients a_i and $N \subseteq \{1, 2, \dots, n+1\}$ the set of negative coefficients of a_i . So we have:

$$\sum_{i \in P} a_i x_i = \sum_{j \in N} |a_j| x_j$$

$\text{VCdim}(\mathcal{HS}_0^n)$

Take $P \subseteq \{1, 2, \dots, n+1\}$ the set of positive coefficients a_i and $N \subseteq \{1, 2, \dots, n+1\}$ the set of negative coefficients of a_i . So we have:

$$\sum_{i \in P} a_i x_i = \sum_{j \in N} |a_j| x_j$$

Assume that A is shattered by \mathcal{HS}_0^n and take $B = \{x_i \mid i \in P\}$. In particular, there exist h_B such that it realizes the label consisting of +1 for all $x_i \in B$ and -1 for all $x_i \notin B$.

$$h_B \left(\sum_{i \in P} a_i x_i \right) = \sum_{i \in P} a_i h_B(x_i) > 0$$

$$h_B \left(\sum_{i \in P} a_i x_i \right) = h_B \left(\sum_{j \in N} |a_j| x_j \right) = \sum_{j \in N} |a_j| h_B(x_j) < 0$$

So, this is a contradiction.

Assignment 1

Assignment 1 – good to know

- 4 problems = 3.5 points
- 1 bonus problem = 1 point
- deadline: in 3 weeks time, Friday, 17th of April 2020, 23:59
 - late submission policy: maximum 3 days allowed, -10% (= 0.35 points) for each day
 - upload a pdf written in a scientific editor (Word, Latex, LyX) containing your solution here: <https://tinyurl.com/AML-2020-ASSIGNMENT1>
 - *is mandatory that you write your solution with a scientific editor, otherwise your solution would not be taken into account*
 - you can insert drawings for your proofs
- for every problem write clear explanations, proofs to justify your answer (if you write just some indications you will not get too many points)
- do not share/copy the solution with/from your colleagues: you + your colleague/s will get 0 points

Problems 1 and 2

Assignment 1

Deadline: Friday, 17th of April

1. (0.5 points) Consider $\mathcal{H} = \{h_\theta: \mathbb{R} \rightarrow \{0,1\}, h_\theta(x) = \mathbf{1}_{[x \geq \theta]}(x), \theta \in \mathbb{R}\} \cup \{h_\theta: \mathbb{R} \rightarrow \{0,1\}, h_\theta(x) = \mathbf{1}_{[x < \theta]}(x), \theta \in \mathbb{R}\}$. Compute $\text{VCdim}(\mathcal{H})$.
2. (0.75 points) Consider \mathcal{H} to be the class of all centered in origin sphere classifiers in the 3D space. A centered in origin sphere classifier in the 3D space is a classifier h_r that assigns the value 1 to a point if and only if it is inside the sphere with radius $r > 0$ and center given by the origin $\mathbf{O}(0,0,0)$.
 - a. show that the class \mathcal{H} can be (ϵ, δ) – PAC learned by giving the algorithm A and determining the sample complexity $m_{\mathcal{H}}(\epsilon, \delta)$ such that the definition of PAC-learnability is satisfied.
 - b. compute $\text{VCdim}(\mathcal{H})$.

Problems 3 and 4

3. (1 point) What is the VC-dimension of the set of subsets I_θ of the real line parameterized by a single parameter θ where $I_\theta = [\theta, \theta + 1] \cup [\theta + 2, \infty)$.
4. (1.25 points) An axis aligned square classifier in the plane is a classifier that assigns the value 1 to a point if and only if it is inside a certain square. Formally, given the real numbers $a_1, a_2, r > 0 \in \mathbb{R}$ we define the classifier $h_{(a_1, a_2, r)}$ by

$$h_{(a_1, a_2, r)}(x_1, x_2) = \begin{cases} 1, & \text{if } a_1 \leq x_1 \leq a_1 + r, a_2 \leq x_2 \leq a_2 + r \\ 0, & \text{otherwise} \end{cases}$$

The class of all axis aligned squares in the plane is defined as

$$\mathcal{H} = \{h_{(a_1, a_2, r)}: \mathbb{R}^2 \rightarrow \{0, 1\} | a_1, a_2, r \in \mathbb{R}, r > 0\}$$

Consider the realizability assumption.

- give a learning algorithm A that return a hypothesis $h_S = A(S)$ consistent with the training set S (h_S has empirical risk 0 on S);
- find the sample complexity $m_H(\epsilon, \delta)$ in order to show that \mathcal{H} is PAC – learnable.
- compute $\text{VCdim}(\mathcal{H})$.

Bonus Problem

Bonus Problem (1 point)

Compute the VC-dimension of the class of convex d -gons (convex polygons with exactly d sides) in the plane. Provide a detailed proof of your result.

$d = 4$

